


Convolutional Neural Networks



Machine Learning Course - PPCIC

About me

Master's degree at CEFET-RJ | Data Architect at TRF2

Admission: May 2018

Conclusion: July 2020

Advisor: Eduardo Bezerra

Project: Apply convolutional neural networks to model spatiotemporal data



rafaela.nascimento@eic.cefet-rj.br



/in/rafaela00castro



/rafaela00castro

Task: image classification

PS: [images are numbers](#) for computers!

Why **not use only** fully connected (FC) layers for this task?

Fully connected layers


- Require a lot of connections >> many weights
- Example: classify grayscale image of size 200x200
- Architecture: suppose only one hidden layer with 500 units
- Neural Net with 20 million weights 🤖
 - Input layer = 40000 units (200 x 200)
 - hidden layer = 500 units

<https://www.cs.toronto.edu/~lczhang/360/lec/w04/convnet.html>

Fully connected layers (2)

Problems of using just FC for image classification:

- Many weights { long processing time
risk of overfitting
- Spatial structure is ignored



<https://towardsdatascience.com/the-most-intuitive-and-easiest-guide-for-convolutional-neural-network-3607be47480>

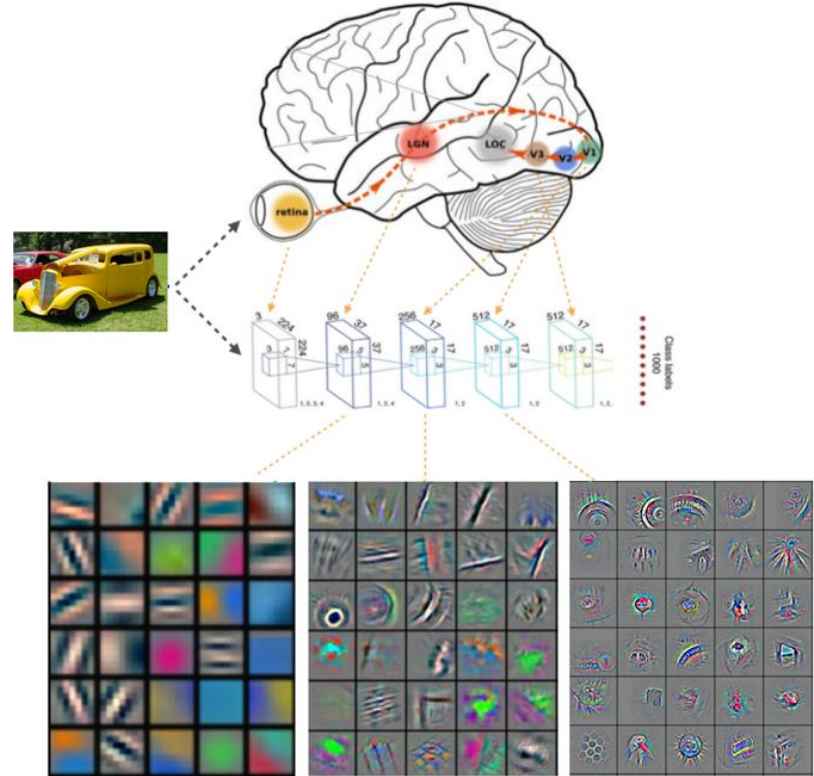
Convolutional Neural
Networks (CNN) are the
solution...

CNN (or ConvNet) introduction

- Introduced by Yann LeCun et al. (1989)
- Inspired in the visual cortex
- Sort images into categories even better than humans in some cases (milestone for deep learning - [AlexNet paper, 2012](#))
- Advantages of using convolutional layers over fully connected layers
 - Local connectivity
 - Weight sharing

Inspired in the visual cortex


- Lower layers detect simple aspects of the image (edges and spots).
- Higher layers detect more complicated patterns (shapes)



<https://medium.com/@vlomonaco/what-i-learned-at-the-deep-learning-summer-school-2017-in-bilbao-c6eae2963554>


Local connectivity (intuition)

- For a particular piece (●), features from the nearby pieces (○) give more information than the pieces further away (○).



Local connectivity

CNN are **not entirely connected**. The neurons in a layer will only be **connected to a small region** (receptive field) of the previous layer.



<https://freecontent.manning.com/deep-learning-for-image-like-data/>

Weight sharing (intuition)

If a **feature detector** knows how to detect a local **feature in one region of the image**, then the same detector is **useful in all other regions** of the image.



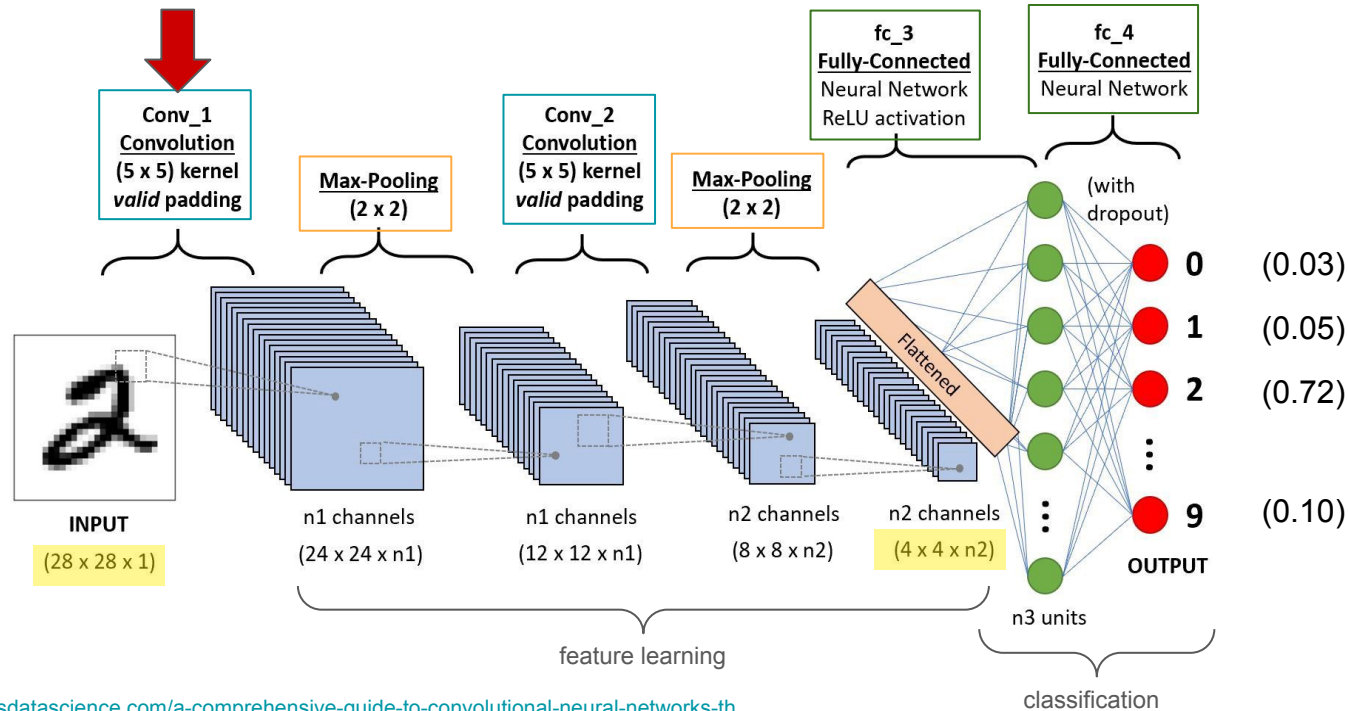
Input

<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

Weight sharing

- In CNN terminology, detector = filter (or kernel).
- The filter consists of a matrix that is learned during training.
- CNN **shares the same parameters** (weights and biases) across different locations in an image (at the same depth slice).

Typical architecture for image classification



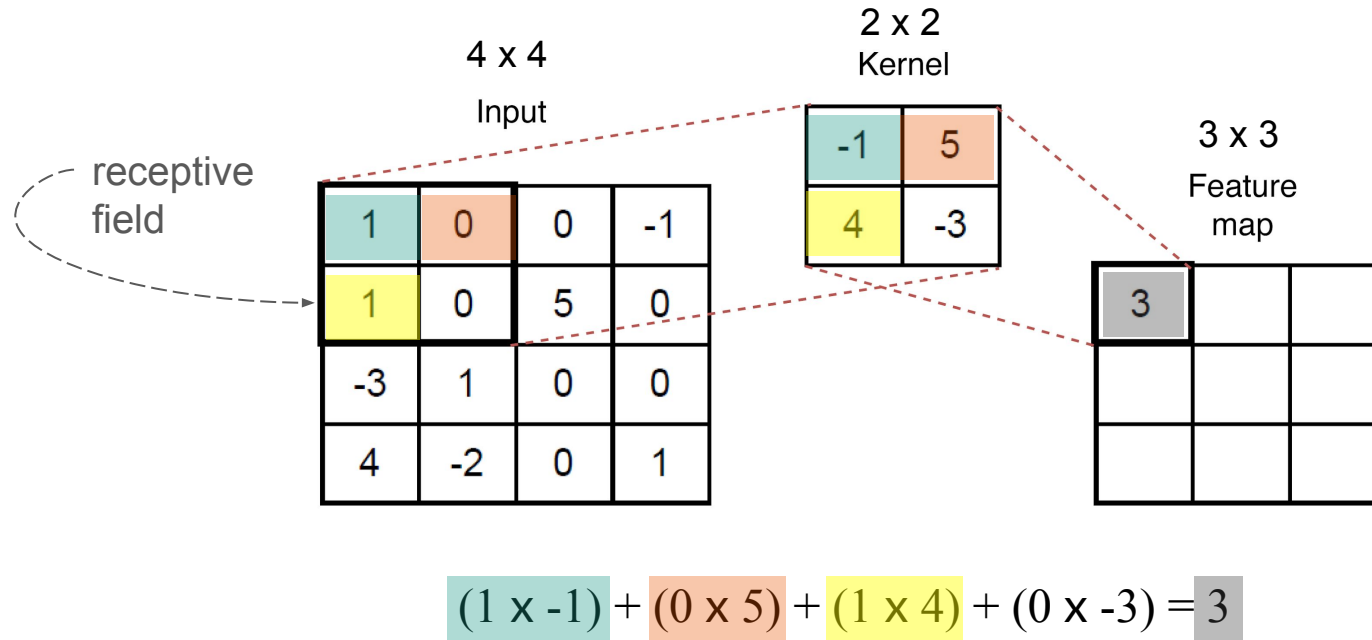
<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Convolution operation

During the **forward pass**:


- The convolutional layer **convolves the filter (or kernel) across the receptive field** of the input volume.
 - i.e., computes dot products between the entries of the filter (weights) and the input at any position.
- The **result** of applying convolution for each filter is a **2-dimensional feature map** (or activation map).

Convolution operation (2)



Convolution operation (3)

The convolutional layer stacks these feature maps (or activation maps) along the depth dimension and produces the output volume.



<https://cs231n.github.io/convolutional-networks/>

Stride

Hyperparameter that controls how the kernel (or filter) slides through the input, i.e., the position at which convolution operation must begin.

(a) Stride = 1

Input grid (6x6):

1	2	3	1	3	5
2	2	5	4	2	5
0	6	9	6	2	2
2	0	1	9	4	0
5	5	4	6	7	6
6	1	3	7	1	5

Kernel (3x3):

1	0	-1
1	0	-1
1	0	-1

Output grid (4x4):

-14	-1	10	-1
-11	-11	7	12
-7	-10	1	13
5	-16	-4	10

(b) Stride = 2

Input grid (6x6):

1	2	3	1	3	5
2	2	5	4	2	5
0	6	9	6	2	2
2	0	1	9	4	0
5	5	4	6	7	6
6	1	3	7	1	5

Kernel (3x3):

1	0	-1
1	0	-1
1	0	-1

Output grid (2x2):

-14	10
-7	1

<https://www.brilliantcode.net/1584/convolutional-neural-networks-1-convolution-layer-stride-padding-kernel/>

Padding

Hyperparameter that defines an additional border in the input. But why?

- Convolution operation shrinks the size of the input >> padding helps generate output equal in size to the input (*same padding*)
- Pixels in the corner are convolved only a few number of times as compared to the central pixels (information loss) >> padding can be used to overcome this issue.

PS: *valid padding* = convolution operation without padding

Padding (2)

2D CNN

stride = 1
padding = 1

zero-padding →

5 x 5
Input size

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

3 x 3
Kernel

0	-1	0
-1	5	-1
0	-1	0

5 x 5
Feature map

114				

<https://www.pyimagesearch.com/2018/12/31/keras-conv2d-and-convolutional-layers/>

Output size calculation

The **dimension** (height x width) of the **output** feature map depends on the **input size** (I), **kernel size** (K), **stride** (S) and **padding** (P).

$$O = \frac{I - K + 2P}{S} + 1$$

$$O = \frac{5 - 3 + 2 \times 1}{1} + 1$$

$$O = 5$$

5 x 5 Input size						
0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

3 x 3 Kernel		
0	-1	0
-1	5	-1
0	-1	0

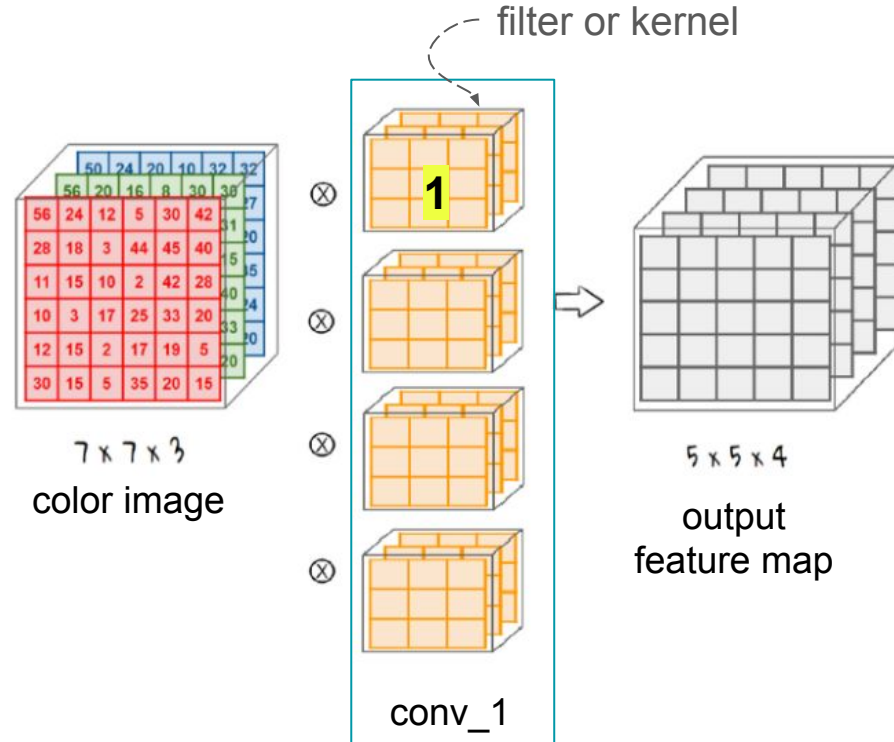
5 x 5 Feature map				
114	328	-26	470	158
53	266			

Example

Given a **color image** that has **7x7 pixel** (this means: the input has a 7x7x3 volume), use one **convolutional layer** with **4 filters** (size 3x3) to detect certain features.


Pay attention to the volume of the output: the **number of filters** will determine the **depth of the output**.

Example (2)



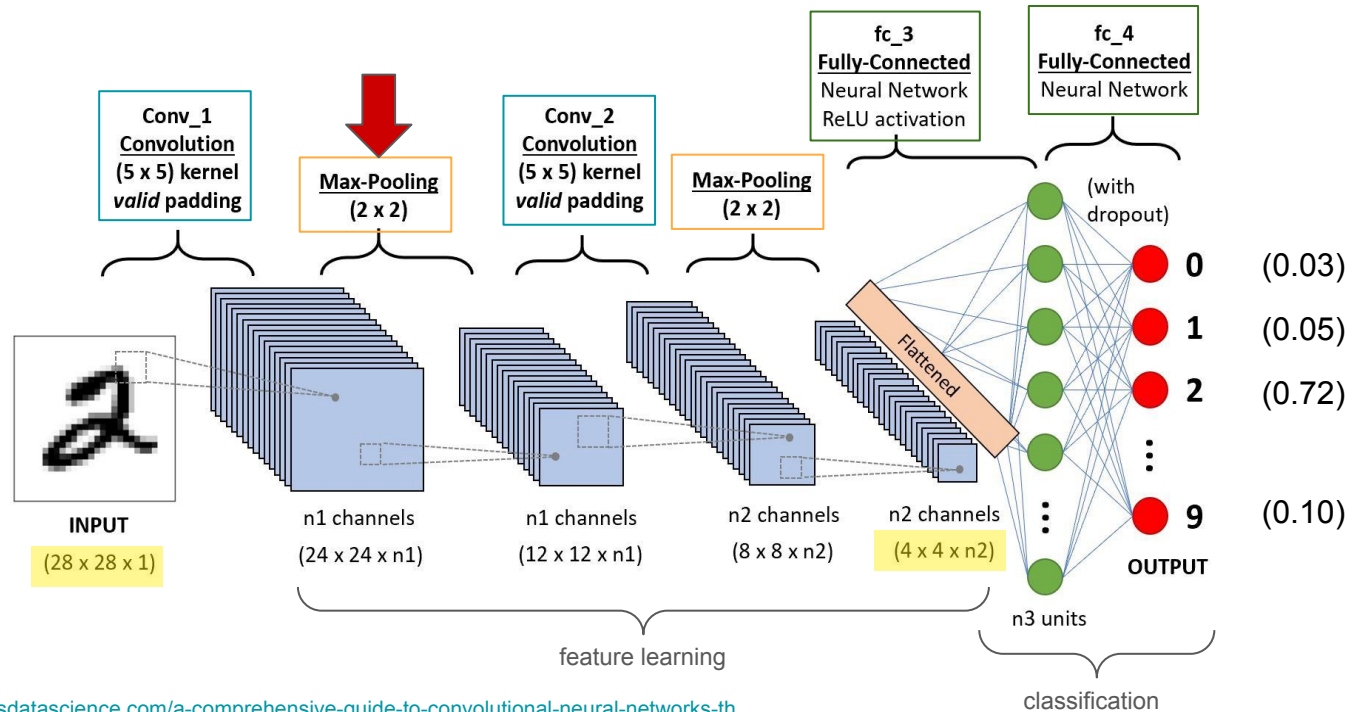
<https://towardsdatascience.com/the-most-intuitive-and-easiest-guide-for-convolutional-neural-network-3607be47480>

Example (3)



<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Typical architecture for image classification





<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Pooling layer

- Commonly used between convolutional layers
- **Reduces spatial dimensions** (width x height) >> decreases the computational power required
- **Extracts dominant features** >> translation invariant
- Depth of the output volume equal to input volume

Pooling layer (2)




<https://cs231n.github.io/convolutional-networks/#convert>

<https://towardsdatascience.com/the-most-intuitive-and-easiest-guide-for-convolutional-neural-network-3607be47480>

Pooling (intuition)

- The three images on the left give the same image on the right.
- The model using only the right image knows that there is the shape on the top left.




<https://www.quora.com/How-exactly-does-max-pooling-create-translation-invariance>

Activation function

Wherever a negative number occurs, swap it out for a 0.

it is easier to train and often achieves better performance.

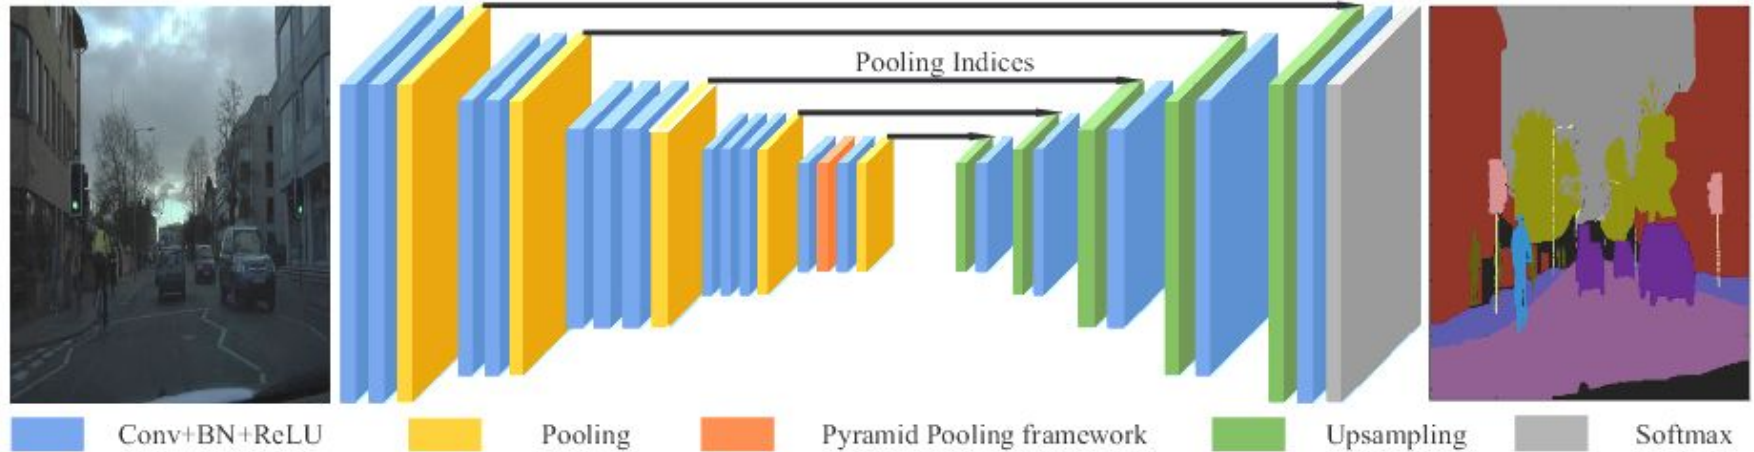
Rectified Linear Unit (ReLU)



<https://towardsdatascience.com/image-classifier-cats-vs-dogs-with-convolutional-neural-networks-cnns-and-google-colabs-4e9af21ae7a8>

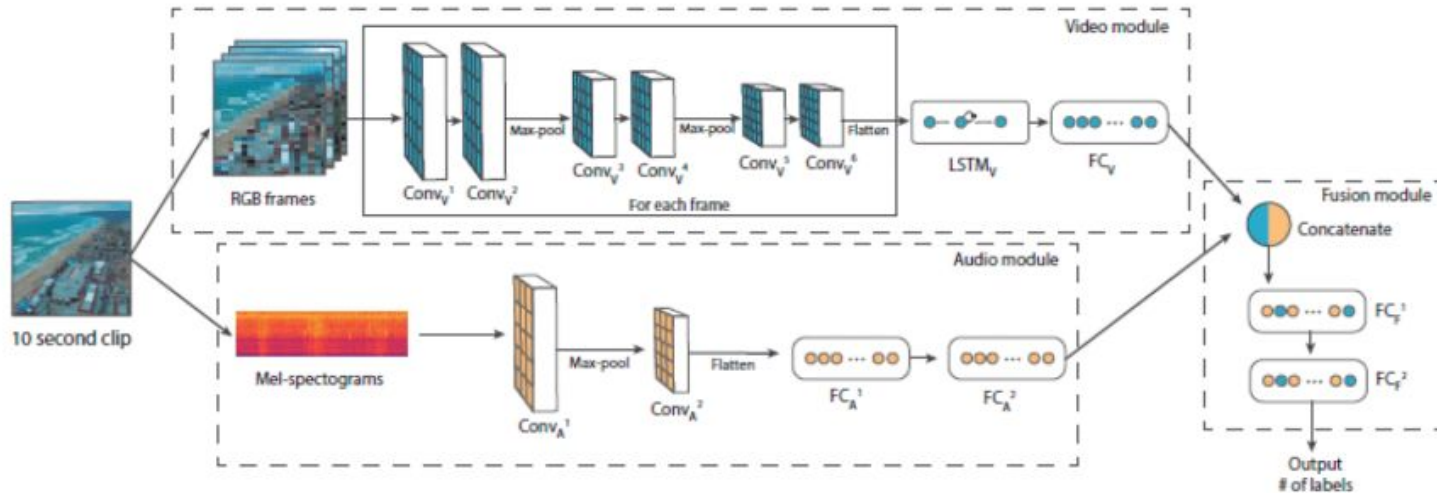
Some applications

Semantic segmentation (2D CNN)



[Tan et al., ICIG 2017]

Effects synchronization (2D CNN)



Bimodal architecture

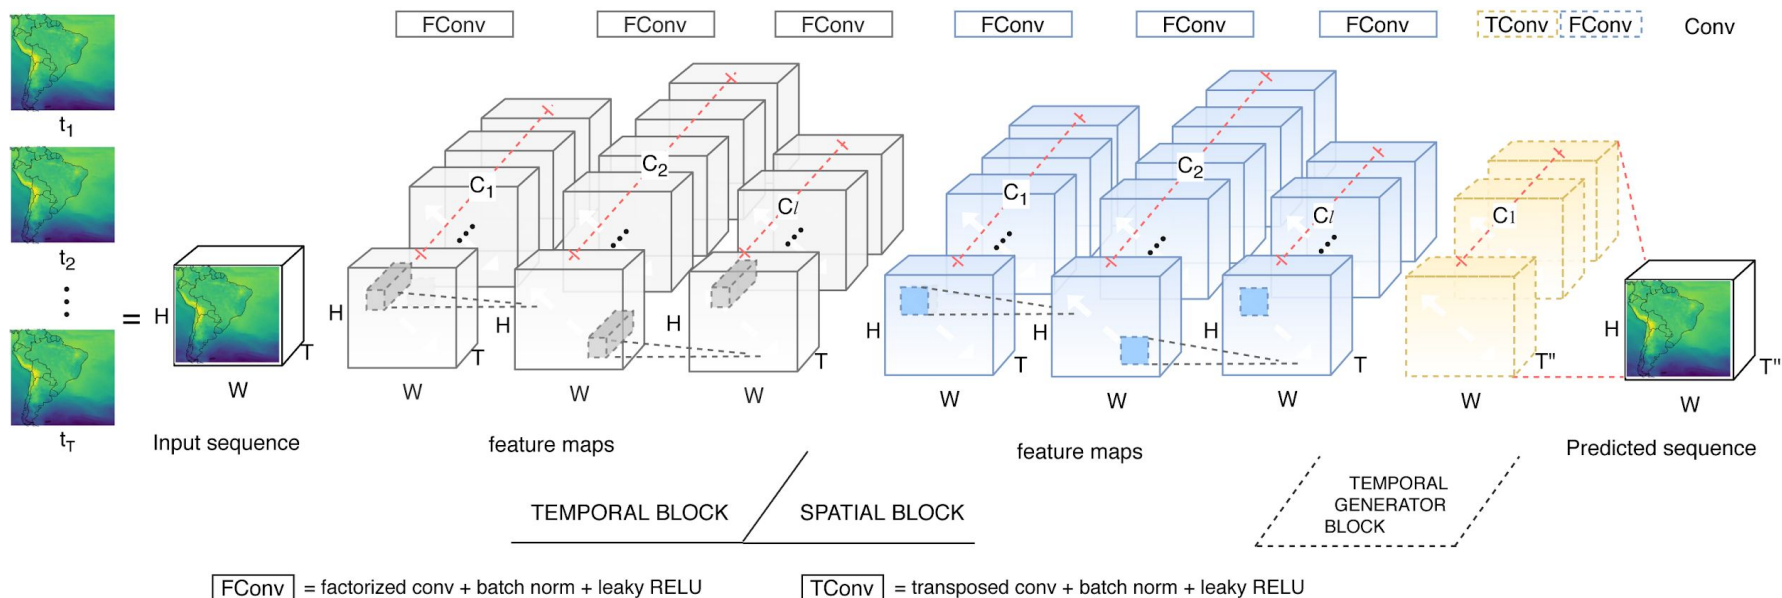
[Abreu et al, IJCNN 2018]

Machine Translation (1D CNN)



[Gehring et al., 2017]

Spatiotemporal data forecasting (3D CNN)



[Castro et al., 2020]


<https://github.com/MLRG-CEFET-RJ/stconvs2s>

Backup slides

Images are numbers



RGB image




[<< back](#)

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

3D CNN

kernel = $3 \times 3 \times 3$
padding = 0
stride = 1



<https://github.com/OValery16/Tutorial-about-3D-convolutional-network/blob/master/images/3dconv.gif>

[<< back](#)